slack_ts	date	poster_name	lr_title	link_to_LR_doc	slack_permalink	summary_of_post	segment	SSv2 (total)	SSv2 Proxy	SSv2 (save/revis/social/download MECE)	SSv2 (intent expressions only)	SSv2 (non-shopping clickout)	SSv2 (shopping clickout)	Repins	Closeups	Pin Searches	Pin Impressions (1px 1s)	Ad Impressions	Web/API Requests (UU)	DAU (All Users)	MAU (All Users)	WAU (All Users)	Novel Domain Impressions (Interests)	Novel Pincept Domains (Interests)	Adopted Pincept Domains (Interests)	other_metrics	metric_source	extras_json
manual:upp-retrieval-clr-v0		Piyush Maheshwari	UPP Retrieval CLR V0	https://docs.google.com/document/d/1AJvmdCNm22AE2hC9056v1GN_cB7hrLQLkP2TCbeOAiU/edit?tab=t.74vo5rr7nv5q		First version of UPP Retrieval CLR on the notifications surface. Stat-sig global WAU gain of 156k with neutral OEC and broad stat-sig lifts across email and push engagement.	All Users																			WAU Global +156k (stat-sig); WAU UCAN +28k (stat-sig); MAU Global +48k (NON-stat-sig); MAU UCAN +28k (NON-stat-sig); OEC neutral; [all stat-sig] Email Unique Non-Unsubscribe Click +0.62%; Push Notification Opened By +0.5%; Email Unique Open +0.14%; Email Click Propensity +0.417%; Push Open Propensity +0.179%; Email Click Volume +1.14%; Email Open Volume +0.375%; Push Open Volume +0.705%; Email Open Rate +0.302%; Email CTR +1.066%; Push Open Rate +0.709%	manual (James-provided; not posted in #p13n-relevance-lr)	{"dropped_non_sig": {"OEC": "stated neutral, no value given", "MAU (All Users)": "Global +48k and UCAN +28k both marked non-stat-sig by the source"}, "notes": "PROVENANCE: added from text James supplied directly on 2026-08-08, NOT from a #p13n-relevance-lr post - the channel has no posts after 2026-07-29. slack_ts is a sentinel string, not a Slack timestamp, so it will not poison the resume cursor; slack_permalink and date are blank because no post exists to derive them from. DOC UNREADABLE: permission-blocked on both the experiments MCP and Drive, and absent from Drive search - so Approval/Ship dates, ship group and Helium link could not be recovered. ALL 18 FIXED COLUMNS BLANK, consistent with the other notification-surface rows (2026-05-18, 2026-05-20): WAU/MAU are absolute counts rather than percentages, and every engagement metric is notification-scoped, none on the canonical 18. SOURCE-TEXT ISSUES worth confirming: (1) \"Push Notification Opened By\" is listed TWICE, as 0.5% and again as 0.50 without a percent sign - recorded once, assuming a duplicate rather than two metrics; (2) UCAN +28k appears as BOTH the WAU and MAU value, which may be a copy error; (3) the source writes \"stag-sig\", read as stat-sig. AUTHOR ATTRIBUTION: the source names Piyush Maheshwari only after the notif_retrieval_upp_v2 block; applied here as same-workstream inference - confirm before citing V0 authorship. LINK MAPPING: the three URLs were assumed to map to V0/V1/v2 in the order given - unverified."}
manual:upp-retrieval-clr-v1		Piyush Maheshwari	UPP Retrieval CLR V1	https://docs.google.com/document/d/1VhOSt9utXdnVUnxRvg10ABSdQuiaKa2IZ1LZy2cTWs/edit		Second version of UPP Retrieval CLR on notifications. Stat-sig global WAU gain of 130k with neutral OEC; email and push engagement lifts are smaller than V0 and fewer of them reach significance.	All Users																			WAU Global +130k (stat-sig); OEC neutral (4.0E-5); Email Unique Non-Unsubscribe Click +0.53% (stat-sig); Push Notification Opened By +0.2% (stat-sig); Email CTR +0.455% (stat-sig); Email Open Rate +0.047% (no significance marked); Push Open Rate +0.12% (no significance marked)	manual (James-provided; not posted in #p13n-relevance-lr)	{"dropped_non_sig": {"OEC": "stated neutral at 4.0E-5", "Email Open Rate": "+0.047% - no significance marker in source", "Push Open Rate": "+0.12% - no significance marker in source"}, "notes": "PROVENANCE: added from text James supplied directly on 2026-08-08, not from a channel post; sentinel slack_ts, blank date and permalink. Doc permission-blocked on experiments MCP and Drive. ALL 18 FIXED COLUMNS BLANK - WAU is an absolute count, all engagement is notification-scoped. COMPARISON WORTH NOTING: V1 reports LOWER gains than V0 on every metric they share (WAU 130k vs 156k, Email Unique Non-Unsub Click 0.53% vs 0.62%, Push Opened By 0.2% vs 0.5%, Email CTR 0.455% vs 1.066%) and reports no MAU at all. The source text does not say whether V1 is measured against the same baseline as V0 or on top of it - that determines whether V1 is a regression or an increment, and it cannot be resolved from the text alone. Author attribution and link mapping are inferred, as for V0."}
manual:notif-retrieval-upp-v2		Piyush Maheshwari	notif_retrieval_upp_v2	https://docs.google.com/document/d/1opgWr2uDkpfP1CARnz8hpAkKnweCihzxxXx_h5gwV4/edit?tab=t.74vo5rr7nv5q		Third iteration of UPP retrieval on notifications. The retention picture inverts versus V0/V1 - MAU is stat-sig (+226k) while WAU is not - and engagement is essentially flat, with the gains instead showing up in content quality.	All Users																			MAU +226k (stat-sig); MAU UCAN +89k (stat-sig); WAU +139k (NON-stat-sig); WAU UCAN +45k (NON-stat-sig); Email Unique Non-Unsubscribe Click +0.15%; Push Notification Opened By -0.01%; Email Open Rate +0.02%; Email CTR +0.00%; Push Open Rate +0.00%; [Content Quality] is_racy -3%; is_low_quality -4.3%; is_non_genai_slop -2.5%	manual (James-provided; not posted in #p13n-relevance-lr)	{"dropped_non_sig": {"WAU (All Users)": "+139k and UCAN +45k both marked non-stat-sig by the source", "engagement metrics": "no significance markers given; Email CTR and Push Open Rate are +0.00%, i.e. flat"}, "notes": "PROVENANCE: added from text James supplied directly on 2026-08-08, not from a channel post; sentinel slack_ts, blank date and permalink. Doc permission-blocked. ALL 18 FIXED COLUMNS BLANK - MAU/WAU are absolute counts, engagement is notification-scoped. THE INTERESTING SHIFT: V0 and V1 both led with stat-sig WAU; here WAU is explicitly NON-stat-sig while MAU becomes stat-sig (+226k), and engagement metrics are flat to slightly negative (Push Opened By -0.01%, Email CTR +0.00%). The case rests on content quality instead. *** FLAG FOR JAMES: \"is_non_genai_slop: -2.5%\" is ambiguous in direction. is_racy -3% and is_low_quality -4.3% are clearly improvements (less racy, less low-quality content). By the same parallel, a DECREASE in \"is_non_genai_slop\" would mean fewer non-slop items, i.e. MORE genai slop - a regression sitting in a list presented as wins. Either the metric name or the sign needs confirming before this row is cited; I have not assumed a direction. *** Author attribution is confirmed for this LR only; link mapping assumed from the order given."}
1785359965	2026-07-29	Yali Bian	HF-LWS Model Scaling	https://docs.google.com/document/d/1fTSxtOsJpK1YSkRFEVgLAOCVv4RBiC8BUtHNknYPMo/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1785359965349809	Scales HF-LWS feature-crossing capacity on GPU via RMSNorm, batched MaskNet and a larger DCN/MaskNet head, sending more relevant candidates downstream to Pinnability. Broad engagement gains at neutral latency and infra cost.	All Users	+0.12%	+0.25%	+0.15%				+0.51%			+0.23%		+0.14%							Board Creates +0.64%; Sessions >5min +0.24%; Pins 28d Fresh Neardup Impressions w/o Ads +0.21%; Torso Non-Fully-Explored Overall TP Repins and Ob5s Clicks (non-ads) -6.49%; Torso Overall Impressions (non-ads) +0.37%; Homefeed Pin Impressions +0.25%; Homefeed Num Pin Click +0.37%; Homefeed Downloads +0.57%; Related Impressions Diluted +0.28%; Related Closeups Diluted +0.28%; Related Repins Diluted +0.50%; Total Time Spent in App (s) +0.21%; Users >15min +0.30%; Users >30min +0.36%; Home Feed (s) +0.27%; Closeup (s) +0.23%; Related Pins (s) +0.23%; Closeup Body (s) +0.23%; Other Surfaces (s) +0.20%	slack (doc blocked)	{"cost":"neutral latency and infra cost","notes":"Closeup (s) +0.23% is timespent, not the Closeups action count - left out of the Closeups column","doc_access":"permission-blocked on both Drive and experiments MCP"}
1784744651	2026-07-22	Ryan Kam	Learned Retrieval Manas Indexing Pipeline Dynamic Field onboarding	https://docs.google.com/document/d/1vlrJBsS0mNBeAAK_4qvhCQdGeNyGCHGL0XdHt6Qm_kQ/edit?usp=sharing	https://pinterest.slack.com/archives/C05UMECTDDJ/p1784744650552249	Eliminates boilerplate per-Learned-Retrieval-model Manas indexing .yaml config PRs by introducing a dynamic field and wiring the indexing pipeline to consume it. Functional no-op vs prod, validated byte-for-byte on resulting metadata.	All Users																			No engagement metrics reported - functional no-op verified by byte-for-byte metadata comparison vs prod and a test-cluster index build/deploy	slack (doc blocked)	{"related_docs": ["https://pdocs.pinadmin.com/docs/manas-user-guide/index_operations/dynamic-field-onboard (Manas dynamic field onboarding guide)"], "notes": "Infra/dev-velocity LR with no experiment arm - all 18 fixed columns legitimately blank. Benefits claimed: dev velocity (no per-model PR) and hardened configuration (centralized shared Manas config). Follow-on to the Unity Experiment Configuration migration LR (2026-07-01).", "doc_access": "permission-blocked on experiments MCP"}
1784675598	2026-07-21	Yali Bian	HF-LWS Training Data Optimization	https://docs.google.com/document/d/1IOk7g6mYZhSytOi1hUNKO-JjH55OG7nYfDfv8eLnaik/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1784675597815519	Deprecates the LWS PROD training dataset in favour of the L2 (Pinnability) sampled dataset as production impressed training data for HF-LWS. Engagement gains plus $80K/year cost saving and reduced upstream operational dependency.	All Users		+0.20%																	Sessions >5min +0.27%; Pins 28d Fresh Neardup Impressions without Ads +0.43%; Homefeed Num Pin Click +0.09% (unique users); Homefeed Downloads +1.14% vol / +0.51% prop; Downloads on Related Pins +0.43%; Pin Reactions on Related Pins +0.70%; Closeup timespent (s) +0.21%; [Core Users] Homefeed Num Pin Click +0.49%; Downloads on HF +1.39%; Pin Reactions on HF +1.74%; Downloads on Related Pins +0.90%	slack (doc blocked)	{"cost": "$80K/year saving", "notes": "\"Closeup (s): +0.21%\" is timespent in seconds, not the Closeups action count - deliberately kept out of the Closeups column. Core User segment values routed to other_metrics and labelled. Foundational benefits claimed: reduced upstream operational dependency, better alignment with online performance, higher dev velocity.", "doc_access": "permission-blocked on experiments MCP"}
1784241609	2026-07-16	J.J Hu	Reflex: Homefeed Following Feed and GraphSage CG Deprecation	https://docs.google.com/document/d/1KWg9o2XUibLtk73OPgDDBqmKlpc2eQcAOGx7wk39ncQ/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1784241608815299	Uses Reflex to deprecate the two weakest candidate generators, Following Feed and GraphSage Product Pin. Engagement and use-case-adoption wins at neutral cost and PWT, with acknowledged shopping and MDD regressions still under discussion.	All Users		+0.17%	+0.12%		+0.28% (unique users)		+0.32%											+0.49%	Unsuccessful SSv2 (total) -0.31% (inverted metric - negative is good); Novel Impressed Pincept Domains (interests) +0.14% (see notes); Homefeed timespend +0.17%; Homefeed Num Pin Click +0.91%; Homefeed Repins +0.73%; Homefeed Clickthroughs +0.56%; Downloads on HF +1.18%; Pin Screenshots from HF +1.84%; Homefeed Use case adoption +0.75%; Homefeed WAU with 2+ use cases +0.54%; unquantified regressions on shopping metrics and MDD, under discussion with content team	slack (doc blocked)	{"cost": "Neutral cost and neutral PWT", "notes": "\"Novel Impressed Pincept Domains (interests)\" is ambiguous between the canonical \"Novel Domain Impressions (Interests)\" and \"Novel Pincept Domains (Interests)\" - routed to other_metrics rather than guessing a column. Doc blocked so it could not be resolved; worth a one-line clarification from J.J. Unsuccessful SSv2 kept out of all SSv2 columns per its inverted definition. Shopping/MDD regressions were disclosed but not quantified in the post.", "doc_access": "permission-blocked on experiments MCP"}
1783548180	2026-07-08	Hedi Xia	Faster Iterative Diversity for L1 Utility	https://docs.google.com/document/d/1llmsQi9kYu9tS4w9wb2nEdLuJdozru7HZRNGoQP-jZU/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1783548180035799	Latency-focused follow-up to the approved Iterative Diversity LR: keeps identical fixed-point diversity behaviour and the three-pass cap but removes a duplicate sort, caches pin utility/cluster assignments, and replaces power operations with a multiplicative multiplier cache. Neutral metrics, -2.5ms p90 and -5.7ms p99.	All Users																			Homefeed latency P90 -2.5ms; Homefeed latency P99 -5.7ms	slack+doc	{"cost": "Neutral infra cost", "dropped_non_sig": {"SSv2": "post and doc both state Neutral, no value given", "Homefeed metrics": "doc states Neutral across the board, no values given"}, "ship_group": "enabled_fp_cap_3_fast", "notes": "Pure latency optimization producing the same ordering up to numerical precision - all 18 fixed columns legitimately blank. Doc carries a numerical-stability trade-off analysis: subnormal/underflow not a practical risk at production settings (smallest diversity factor d>=0.8, count boundary ~2000, giving ~1.5e-194 vs Double.MIN_NORMAL ~2.2e-308). No 12wk WAU impact to report.", "related_docs": ["https://docs.google.com/document/d/10crKwJ-sapsPKVRAAswcdD_LCOlUoTvwb1Ve6jCApc ([LR] Iterative Diversity for L1 Utility - previously approved parent)"]}
1783447874	2026-07-07	Yuantao Ji	Closeup Cache Invalidation for NLFUs	https://docs.google.com/document/d/1r13c9RRMPAWvUWPV5K0_ifGtHFMgevtmychxMsx2TA/edit	https://pinterest.slack.com/archives/C05UMECTDDJ/p1783447873989479	Adds PIN_CLOSEUP as a Homefeed chunk invalidation trigger for non-core non-casual users, so the next Homefeed page fetches a regenerated chunk. Delivers one of the largest NLF WAU wins to date, recommended to ship on the legacy user context sequence ahead of the pending real-time sequence.	NLFUs (non-core, non-casual users)		+0.28%									+0.33%								NLF WAU +0.037% diluted (37k WAUs) / +0.0472% undiluted, p=0.03; HF Repins +0.33% prop; HF Num Pin Click +1.12%; HF Long Clickthroughs +1.78% vol / +1.04% prop; Downloads +0.87% vol / +0.39% prop; Screenshots +1.96% vol / +0.85% prop; Pin Share on HF +1.75% vol / +1.59% prop; HF Pin See More +0.89% prop; unity-homefeed blending_latency +4.99%; HomefeedAllStructuredFeedPlanner_overall_latency +1.66%; HF Pin Reports propensity +18.9% ESC-5 and US users, +16.8% US users; P1 Actioned User Report 24.081%; P2 Inclusive Filter Adoption Body Type -10.7%	slack+doc	{"helium_link": "https://helium.pinadmin.com/hf_chunk_invalidation_lfu/combined/metrics-v3/ (experiment hf_chunk_invalidation_lfu, version 1, control/enabled, snapshot 2026-06-07)", "cost": "+$63,512/year pre-optimization; post proposes the cost be absorbed by a unified NLFU V-team budget", "notes": "Fixed columns are for the NLFU segment, not All Users - this is the only row in the tracker so far that is not All Users. Ad Impressions +0.33% came from the doc only; the Slack post omits it. NLF WAU is a cohort-qualified retention metric and is kept out of the WAU (All Users) column per the traps list. Post cites diluted +0.037%, doc Summary cites undiluted +0.0472% - different measures rather than a contradiction, both recorded. Doc guardrail row flags P1 Actioned User Report at 24.081% and a Trade-offs section on HF Pin Reports propensity, neither surfaced in the Slack post. Recommendation is explicitly an interim ship on the legacy UCS; no firm timeline on the real-time L30 sequence that would supersede it. Approval Date and Ship Date blank in doc.", "related_docs": ["https://docs.google.com/document/d/1UZhgv0kuNLimF59pZv37ortB8lORX49MLWEEtrYso0M (Experiment Proposal: closeup/screenshot/save-based cache invalidation for resurrected users)", "https://docs.google.com/presentation/d/1tAfqKgxqAwuoLm8VFfjH91SApjSjCKgXf18v2bz3mkA (In-Session Responsiveness deck)", "https://docs.google.com/document/d/1V6PdZr_QKmT06blKzj60zPWBc4UBLgY90Za9o8Fo4iU (In-session responsiveness strategy on Homefeed)"]}
1782851753	2026-06-30	Hanlin Lu	HF-RecGPT replaces Multi Embedding Sizer	https://docs.google.com/document/d/1up0tCcNiM_bKTPDtDHuxaBwFxbz3GQMeKSJ3hTy3JhY/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1782851752999499	Reallocates 150 sizers from Multi Embedding to HF-RecGPT, roughly doubling RecGPT pin coverage and its impression share. Sitewide repin and SSv2 MECE gains at neutral cost, with a small intent-expressions regression.	All Users			+0.18%	-0.13%			+0.42%												Save sessions +0.25%; Revisitation sessions +0.26%; Download sessions +0.16%; SSv2 save/revis/social/download ONLY (non-MECE) +0.17%; SSv2 save/revis/social/download OR clickout +0.15%; Homefeed repins +0.64% vol / +0.26% unique users; HF-RecGPT pin coverage +100%; user coverage +20%; impression share +89.3%	slack (doc blocked)	{"cost": "Neutral cost changes", "notes": "Budget-reallocation LR: gains come at Multi Embedding's expense, so the sizer trade is the decision, not the lift. Post reports BOTH \"save, revisitation, social, and download (MECE)\" +0.18% and \"save, revisitation, social, and download only\" +0.17% as separate lines - only the explicit MECE value fills column 10, the \"only\" variant goes to other_metrics per the traps list. Intent expressions -0.13% is a genuine stat-sig regression on a canonical goal metric.", "doc_access": "permission-blocked on experiments MCP"}
1782410368	2026-06-25	Ryan Kam	Learned Retrieval Model Config Migration to Unity Experiment Configuration	https://docs.google.com/document/d/1nSdXnNfZiQhaFpQ4CLy88TlpVlkYuFhrg81-mwbr3SM/edit?usp=sharing	https://pinterest.slack.com/archives/C05UMECTDDJ/p1782410368379009	Migrates Learned Retrieval model configs from PinAdmin ManagedData to statically typed Unity Experiment Configuration in the PinConf repo, putting config changes through the normal code lifecycle. Reduces oncall load and hardens a freeform-text configuration layer.	All Users																			No engagement metrics reported - config-migration LR. Claimed benefits: reduced oncall load (targets "Model not found rate is too high", perf_test failures, OOM inference job failures), code-reviewed config changes, AI-agent-editable configs, elimination of freeform-text config errors	slack (doc blocked)	{"related_docs": ["https://docs.google.com/document/d/1t6GwkJceYH-ulu2pW2FCBc8S8odZb6DNWLIqZYGmsfw (streamline Manas index-building config by eliminating redundant .yaml PRs - the follow-on that shipped 2026-07-22)"], "notes": "Parent of the 2026-07-22 Manas dynamic-field LR by the same author; this pair is one workstream. All 18 fixed columns legitimately blank - no experiment arm. Drew 16 thread replies, the most discussion of any June LR.", "doc_access": "permission-blocked on experiments MCP"}
1782229306	2026-06-23	Andreanne Lemay (She Her)	Personalized Interest Exploration in Blending	https://docs.google.com/document/d/1Z3OkQa31Nwg0RthCfDs5SicUxtl-_soMRSswAv4jzz8/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1782229305520809	Adds interest exploration to the blending exploration node, promoting one pin per request to a fixed slot and balancing not-yet-graduated pins (new to the app) against interests new to the user. Engagement, use-case-adoption and WAU gains at the cost of more hides and reduced fresh gains.	All Users		+0.23%																+0.38%	Closeup timespent (s) +0.25%; Sessions >5min +0.28%; HF pin impressions +0.25%; HF pin shares completed +0.73% prop; HF See More +3.59%; Novel Impressed Pincept Domain (interests) +0.25% (see notes); Time spent on app +0.22% (HF +0.29%, closeups +0.25%, other profiles +1.07%); SSv2 Social Sessions only +0.42%; SSv2 Save Sessions only +0.24%; WAU +0.08% to +0.09% stat-sig 6/14-6/17; Low frequency user state v2 WAU stat-sig gain; TRADE-OFF Pin hides +2.90%; reduced fresh/non-grad gains vs hf_content_exploration	slack+doc	{"helium_link": "experiment hf_interest_exploration_fl", "cost": "No additional infra cost - node and signal already exist. Neutral PWT.", "ship_group": "enabled_first_eligible_slot_5_imp_10_maxidx_200 (~35% of requests get interest exploration, 65% non-grad)", "notes": "\"Novel Impressed Pincept Domain (interests)\" is the same ambiguous label seen in J.J Hu's 2026-07-16 LR - sits between canonical Novel Domain Impressions (Interests) and Novel Pincept Domains (Interests). The doc repeats the post's wording verbatim and does not disambiguate, so it stays in other_metrics. WAU gain is date-windowed (6/14-6/17) rather than full-period, so it is NOT put in the WAU (All Users) column. Doc quantifies the hide increase as localized to slot 5. This LR explicitly cannibalizes the fresh gains of the earlier hf_content_exploration LR - a deliberate exploration-budget reallocation.", "related_docs": ["https://docs.google.com/document/d/1iWZERcAWGn0p_ExnsdHs-DMTE043wAYqXNl5OVn6TWc ([LR] hf_content_exploration - the parent exploration node)"]}
1781198598	2026-06-11	Hedi Xia	Iterative Diversity for L1 Utility	https://docs.google.com/document/d/10crKwJ-sapsPKVRAAswcdD_LCOlUoTvwb1Ve6jCApc/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1781198598004049	Reframes Homefeed diversity from finding a heuristic optimum of an NP-hard problem to running an optimizer that incrementally improves orderings over three rerank passes. Engagement and diversity gains for one added for-loop, at the cost of +3.84ms latency.	All Users		+0.17%					+0.32%			+0.15%									Homefeed impression +0.21%; Use Case Intention +0.14%; Impression Diversity +0.16%; Homefeed use case adoption +0.23%; [Core Users] SSv2 Proxy +0.25%; Pin Impression +0.26%; Homefeed impression +0.33%; TRADE-OFF latency +3.84ms	slack+doc	{"cost": "Neutral infra cost; latency +3.84ms accepted as recoverable overhead", "ship_group": "enabled_fp_cap_3", "notes": "Parent of the 2026-07-08 Faster Iterative Diversity LR, which recovered the latency to -0.3ms vs production. Post pre-announces that follow-up, so the +3.84ms was presented as temporary by design. Core Users segment values routed to other_metrics and labelled. Doc frames the latency as \"overhead, not intrinsic to iterating - so it is recoverable\", and names an L3 Diffusion Reranker as the P1 next step.", "related_docs": ["https://docs.google.com/document/d/1gE9I03Wj6SpPlm-ARBGTZcXkt0jvBxp3B0nWzev_LW4 ([LR] Retentive Recs: UIC & SID for Diversity Control on L1 Utility)", "https://docs.google.com/document/d/1Kqr1LCMrYj0L3hYTAYCJDoihjYGU2-8i3x9CU5VgbK8 (Fixed-Point Diversity for L1 Utility)"]}
1780951119	2026-06-08	Kyle Soares	HF + BMI RecGPT UserEventsView Migration (2 lightweight LRs)	https://docs.google.com/document/d/1wSvCF3ME927AuB9DV3P1DPpf2LoCvRNDU-s4K4B8PkA/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1780951119193459	Migrates HF and BMI RecGPT off a custom legacy user-sequence serving stack onto the standardized UserEventsView infra. Neutral metrics at zero cost, unifying serving paths and unblocking the RecGPT SMS migration.	All Users																			[HF] topline neutral, homefeed overall neutral, long clickthroughs +0.66% actions/user; [BMI] neutral on surface, non-shopping clickout SSv2 -0.39% actions/user and -0.24% unique users (team attributes to pre-AA effect, ad revenue neutral, accepted as trade-off); user sequence differences in 7/1000 user_ids pre-UEV vs UEV from real-time signal flakiness and serving-time enrichment	slack (doc blocked)	{"related_docs": ["https://docs.google.com/document/d/1TdSjPNto5tZj95rS8C-j0HcZLI6oJRKRtFu5gCrDCRk ([LR] BMI RecGPT UserEventsView Migration - the second LR in this post)"], "cost": "$0 - GSS QPS and CG budget unchanged", "dropped_non_sig": {"SSv2 (non-shopping clickout)": "BMI -0.39% left out of the column: team attributes it to a pre-AA effect with neutral ad revenue, and Matt Chun confirmed it as an acceptable trade-off"}, "notes": "ONE row covering TWO LRs announced in a single post (HF and BMI), per the skill's single-doc rule - link_to_LR_doc is the HF one, BMI is in related_docs. Sheet import may warrant splitting into two rows. Unblocks RecGPT Triton to SMS migration, in-session user sequences, and the OmniSage upgrade.", "doc_access": "permission-blocked on experiments MCP"}
1780676925	2026-06-05	Kelly He	Retentive Recs x DR: New Use Case Expansion V2	https://docs.google.com/document/d/1bME_hfkXbY5NVRN9memPPCOUuXNaFaZ9Nlfyw1RtsiP4/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1780676924633739	Expands new-use-case rewards in Downstream Rewards, delivering DAU/WAU and SSv2/SRSD lifts concentrated in non-core and casual users, plus increased pin novelty. Includes an infra change logging new-use-case labels for reuse beyond Pinnability.	All Users																			[All Users] DAU +0.092%; WAU +0.066%; Total Impressions Volume on Novel Domains (Interests) +0.08% prop; Novel Impressed Pincept Domains (Interests) +0.08% prop; [Non-core/casual users] DAU +0.229%; WAU +0.141%; SSv2 total +0.31%; SRSD +0.46%; Total Impressions Volume on Novel Domains (Interests) +0.12% prop; Novel Impressed Pincept Domains (Interests) +0.12% prop	slack (doc blocked)	{"notes": "DAU +0.092% and WAU +0.066% are All-Users values and DO belong in the DAU/WAU columns on a strict reading - deliberately left in other_metrics instead because the post leads with the non-core/casual segment and presents All Users as the secondary cut, making the segment attribution of the headline claim ambiguous. Worth one clarifying question to Kelly before the sheet import, since these would be the first populated DAU/WAU cells in the tracker. \"Novel Impressed Pincept Domains (Interests)\" is the recurring ambiguous label - third occurrence in this tracker. Infra change logs new-use-case labels in DR/Ray for reuse by CLR, closeup and search.", "doc_access": "permission-blocked on experiments MCP"}
1780608694	2026-06-04	J.J Hu	Migrate Shopping L2R CGs to CLR	https://docs.google.com/document/d/1wp8Y8nPpYwIOEbr6uASzk0B-A--aXFNDO53bXMBJQ8/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1780608693645509	Migrates Shopping L2R from a stale legacy model to CLR on the latest training data, served with an empty condition across all three shopping corpora. Notable as the first end-to-end production experiment surfaced and implemented through the Reflex agent stack under engineer supervision.	All Users																			No metric values reported in the post - framed as foundational modernization. Claims: retires stale legacy model and its maintenance burden; partially addresses the 2-3x relevance gap vs organic CGs surfaced by Reflex Detect agents	slack (doc blocked)	{"notes": "Reflex milestone LR: opportunity detected autonomously by Reflex Detect agents (PM + DS), implementation generated by the Reflex Build Agent under engineer supervision. Strategically the most significant June LR for the Reflex workstream even though it reports no metrics. Post explicitly separates the follow-up (deprecate GraphSage-shopping CG and re-tune shopping CG budgets) to isolate the budget-reallocation question - that follow-up landed as the 2026-07-16 Reflex CG deprecation LR. All 18 fixed columns blank: no metrics stated.", "doc_access": "permission-blocked on experiments MCP"}
1780533136	2026-06-04	Tuan Truong	PinnerSpark V0.1 Upgrade in Homefeed Conditional Learned Retrieval	https://docs.google.com/document/d/1Q7H1bk5AaKK0VVkktQG6n6HBUfZyMsPCcVcRUAEatTs/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1780533136484609	Upgrades the PinnerSpark interest signal used in Homefeed CLR with a better LLM, richer pin representation and a fix removing stale v0 signals. Broad engagement and use-case gains, markedly stronger for P10 users, with infra savings from the stale-signal cleanup.	All Users		+0.15%								+0.12%	+0.31%								Sessions >5min +0.16%; HF Pin Impressions +0.15%; HF Repins +0.14%; HF Interest Grid Clicks / Closeups +1.17%; HF See More +0.81%; Use Case Intention +0.13%; Use Case Retention +0.27%; [P10 Users] SSv2 Proxy +0.52%; Total Successful Sessions +0.23%; HF Pin Impressions +0.65%; HF Num Pin Click +0.52%; HF Repins +1.20%; Downloads on HF +1.26%; Sessions >5min +0.53%; Retention WAU +0.34%; Retention MAU +0.31%; TRADE-OFF Trustworthy Product <5s Short Clicks +0.60% APU / +0.29% UU; pin-flag regression in subtype analysis (doc only)	slack+doc	{"cost": "Directionally favorable from stale-signal cleanup: -$28K/year stat-sig-only conservative estimate, -$250K/year including non-stat-sig services. No PWT/latency impact.", "ship_group": "enabled_r3_p3_shuffle_20_global_ordered", "conflicts": {"Ad Impressions": {"slack": "not reported in the post", "doc": "+0.31% in Actions/Users"}}, "notes": "Doc adds Ad Impressions +0.31%, absent from the post. Doc also discloses a pin-flag regression in the subtype analysis called out as a localized watchout for post-launch monitoring - the post only says aggregate Pin Reports/Flags are neutral, which is true but less complete. P10 WAU/MAU are segment values and stay out of the All-Users WAU/MAU columns. Doc hypothesizes the strong P10 result comes from v0.1 better handling non-en languages. James is listed as both experiment design reviewer and launch reviewer.", "related_docs": ["(Onsite) PinnerSpark Deep Dive (linked in doc)"]}
1780360529	2026-06-02	Sujan Reddy Ale	UPP: Foundation Model in CLR	https://docs.google.com/document/d/1BD62SWZ_1-H7aChzmlBtRDOZN_6B2Rfxl2toqv6rg/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1780360528626209	Integrates the Foundation Model into the unified Homefeed CLR model, lifting Homefeed repins and the repin rates of every CLR candidate generator. Crucial for scaling the UPP base model, but comes with download and freshness regressions.	All Users	+0.07%	+0.20%	+0.15%				+0.49%	+0.22%	+0.22%										5s+ Outbound Clicks +0.34%; Sessions >5min +0.16%; [Homefeed] repins +0.85%; clickthroughs +0.45%; Screenshots +2.09%; Pin Hides -3.65%; Pin Reports -6.86%; Interest Grid Clicks (Closeups) +1.15%; [CG repin rate] Board CLR +7%; Rec. Topics Interest CLR +3.5%; UIC CLR +3.4%; Pin CLR +2.4%; TRADE-OFFS Homefeed Downloads -1.53%; Fresh impressions -1.73%; Fresh repins -3.17%	slack+doc	{"ship_group": "enabled_fm_cs_and_closeups_15", "notes": "\"SSV2 (SRSD) +0.15%\" mapped to the MECE column on the strength of the Downstream Rewards doc (2026-06-16), which spells the same shorthand out as \"successful sessions v2 (srsd mece only)\" - a cross-LR inference, flagged here rather than silent. Pin Searches +0.22% is the ONLY populated Pin Searches cell in the tracker; the post writes it as \"Pin searches*: 0.22%\" with a stray asterisk and no sign, read as positive from context. Fresh regressions (-1.73% impressions, -3.17% repins) and Downloads -1.53% are the real decision content: doc notes an SRSD-head variant could mitigate downloads and defers it to a future iteration. Launch reviewer is Dylan only.", "related_docs": ["https://docs.google.com/document/d/1cVq3Em_-0bNcV7X-h3o9QHsQjv2Xg1hQ3elh9zh4Uc (CLR FM)"]}
1780068356	2026-05-29	Devin Kreuzer	CLR Condition Budget Tuning via Claude Integration	https://docs.google.com/document/d/182mhkGPRdEyaBsaxYyZXePrWWlaKlhAuFFLi7rNljk/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1780068356468789	Uses Claude to automatically tune CLR condition budgets from experiment feedback, exploiting the serving capacity freed by the GPU Serving launch. Strong use-case and diversity wins with SSv2 gains, traded against freshness regressions and a small cost increase.	All Users		+0.18%	+0.09%						+0.19%	+0.21%		+0.11%						+0.32%	Impression Diversity +0.44%; Homefeed Use Case Adoption +0.14%; Total Impressions Volume on Novel Domains (Interests) +0.19%; Novel Impressed Pincept Domains (Interests) +0.44% (see notes); Adopted Pincept Domains (Interests) Users +0.06%; Board Creates +0.49%; Sessions >5min +0.15%; [Homefeed] Impressions +0.38%; Closeups +0.27%; Downloads +0.48%; Pin Reactions +1.06%; Pin Shares +0.83%; Timespent +0.38%; Longclickthroughs -0.20%; TRADE-OFFS Fresh Impressions without Ads -0.20%; Fresh Long Clicks without Ads -1.24%	slack (doc blocked)	{"cost": "+~$26K from increased Manas queries; SMS cost neutral", "notes": "AI-in-the-loop LR: Claude used to tune condition budgets via experiment feedback - methodologically notable alongside the Reflex agent-stack LR (2026-06-04). \"SRSD +0.09%\" mapped to the MECE column on the same cross-LR basis as other SRSD shorthand (see the 2026-06-16 Downstream Rewards doc), flagged rather than silent. \"Novel Impressed Pincept Domains (Interests) +0.44%\" is the recurring ambiguous label - 4th occurrence in the tracker. Enabled by the GPU Serving launch, so this is downstream of a capacity win.", "doc_access": "permission-blocked on experiments MCP"}
1779409650	2026-05-22	Yali Bian	LWS Funnel Log Removal (unimpressed data as early-funnel negatives)	https://docs.google.com/document/d/1D50Yg2ipY4HIlRLrT5SEpCtZcycmaDGZhIL2ebhKIFA/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1779409650194399	Removes funnel log from LWS training, relying on the broader unimpressed distribution as the early-funnel negative source. Engagement gains plus $46K/year saving from sunsetting funnel logging and its downstream workflow and S3 storage.	All Users	+0.10%	+0.15% vol / +0.04% prop	+0.10%																[Homefeed] Pin Impressions +0.23%; Num Pin Click +0.31%; Clickthroughs +0.64%; Long Clickthroughs +1.04%; Downloads +1.03% vol / +0.37% prop; [Closeup] Related Impressions Diluted +0.14%; Related Closeups Diluted +0.21%	slack (doc blocked)	{"cost": "$46K/year saving from sunsetting funnel logging, downstream workflow processing and S3 storage", "notes": "MECE value written \"+0.10\" without a % in the post - recorded as +0.10% since every sibling value is a percentage. Predecessor in the same LWS training-data line as the 2026-07-21 HF-LWS Training Data Optimization and 2026-07-24 Position-Aware Sampling LRs. Post frames this as a foundation for 16k logging and teacher distillation.", "related_docs": ["https://docs.google.com/document/d/1CstdMibmwK0ci-JjyR0XqUSPm4_H71kPEh70DO3VbR0 (previous LR establishing the principles)", "https://docs.google.com/document/d/1mHEpVlt7j2i7agksSn1-n7Gvx10iGpy3LrJeBw0cMDo (16k logging)"], "doc_access": "permission-blocked on experiments MCP"}
1779395518	2026-05-21	J.J Hu	Sub-Chunk Request CG (Homefeed chunk cache recycling)	https://docs.google.com/document/d/12o-3p-UQmcAs8hDFMyL5RFdnouadhvZSINo8SS2XlY/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1779395517757559	Onboards a CG that recycles unconsumed pins by retrieving directly from the Homefeed chunk cache on 100% of requests, skipping both LWS and Pinnability for a cost-negative profile. SSv2, retention and use-case gains plus a $48K/year saving.	All Users	+0.14%		+0.20%																SSv2 (shopping CLOSEUP sessions only) +0.49% (see notes - not the canonical shopping clickout metric); SSv2 total first-12-days +0.19% stat-sig, full-period trending significant with latest-day CI [-0, 0.0216]; [UCAN] SSv2 Total +0.29% (close to significant); SSv2 intent expressions only +0.51%; SSv2 non-shopping clickout +0.95%; Sessions >5m +0.29%; Total Time Spent in App (s) +0.25%; Users >15min +0.32%; Use Case Intention +0.24%; Homefeed Use Case Adoption +0.37%; Use Case Retention +0.35%; Impression Diversity +0.13%	slack (doc blocked)	{"helium_link": "https://helium.pinadmin.com/hf_sub_chunk_request_cg_exp/combined/metrics-v3 (version 6, control_1|score_sort_0_100)", "cost": "$48K/year SAVING, versus a ~$150K/year cost INCREASE in the prior experiment - the cost profile is the headline change", "notes": "\"SSv2 (shopping closeup sessions only) +0.49%\" is NOT the canonical \"shopping clickout\" metric - closeup and clickout are different actions, so it stays in other_metrics. SSv2 total +0.14% recorded with the caveat that the post itself describes it as only \"trending significant\" over the full period with a CI touching zero; the clean stat-sig read is the first-12-days +0.19%. UCAN segment values routed to other_metrics. Architecturally notable: skips LWS and Pinnability entirely, and the post says it unlocks refocusing other CGs on in-session signals.", "doc_access": "permission-blocked on experiments MCP"}
1779306690	2026-05-20	Jimmy Lin	homefeed_neardup_b_flink_experiment	https://docs.google.com/document/d/1go3KFmQmPBlKrYBPB6WsUJDlVOdWFGDM8baFMuzsv2k/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1779306690436629	Migrates the HOMEFEED_NEARDUP_B Aperture segment from Pinpin indexing to the Flink-powered HOMEFEED_NEARDUP_B_FLINK segment indexed by ImpressionTracking. Run purely to confirm the new indexing job introduces no regressions; result was neutral as expected.	All Users																			Metric result neutral as expected - no values given	slack (doc blocked)	{"dropped_non_sig": {"all metrics": "post states \"The metric result is neutral which is expected\" with no values given"}, "notes": "Pure infra-parity LR - all 18 fixed columns legitimately blank. Shortest LR post in the tracker.", "doc_access": "permission-blocked on experiments MCP"}
1778691704	2026-05-13	Yichu Zhou (He Him)	Homefeed Shopping Load Tuning with Deep RL	https://docs.google.com/document/d/1jaToXFZko8l05G2wIGP74NB_dqViTVrYW1D0S_rAdEo/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1778691704025289	Replaces the tabular RL policy for request-time CG shopping load tuning with a neural network policy learning from richer user- and request-level signals. Sitewide SSv2 and engagement gains at near-neutral infra cost and GPS; P6 countries only.	All Users	+0.15%	+0.18%	+0.17%	+0.19%	+0.34%														Sessions >5min +0.21%; Homefeed Num Pin Click +0.31%	slack (doc blocked)	{"cost": "Infra cost and GPS nearly neutral", "notes": "Scope caveat: change only affects P6 countries, so these are not global-population effects despite the All Users label - the post is explicit about this. \"SRSD session +0.17%\" mapped to the MECE column on the cross-LR basis established by the 2026-06-16 Downstream Rewards doc. Second-generation RL policy replacing the earlier tabular one.", "related_docs": ["https://docs.google.com/document/d/11Ebf7nB8m4PTz-YJiqb_YZ1bal5-5vONQdVGFL_59TI (earlier tabular RL policy LR)"], "doc_access": "permission-blocked on experiments MCP"}
1778094843	2026-05-06	Hanlin Lu	HF-RecGPT diversity improvements	https://docs.google.com/document/d/1ne87dxoTm22K1DOq6zLyTNzBWg2_4nqunQmDjwsJpY/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1778094842675639	Adds Future Window Token Prediction and Gaussian Noise Application to HF-RecGPT to improve diversity and recall. Topline and Homefeed gains with a 50% pin coverage boost, neutral cost on the newly approved Manas cluster.	All Users	+0.08%	+0.16%	+0.10%																[Homefeed] Impressions +0.22%; Pin Click +0.28%; Repins +0.38% vol / +0.15% prop; Clickthroughs +0.23% vol / +0.21% prop; Downloads +0.55%; Screenshots +0.66% vol / +0.36% prop; HF-RecGPT pin coverage +50%; impressions +41.1%	slack (doc blocked)	{"cost": "Neutral when launching on the newly approved Manas cluster", "notes": "Same author and workstream as the 2026-06-30 HF-RecGPT replaces Multi Embedding Sizer LR - this is the model-quality step, that one the budget-reallocation step. \"SSv2 SRSD +0.10%\" mapped to MECE on the cross-LR basis. Coverage gains (+50% pin, +41.1% impressions) are the mechanism behind the later sizer reallocation case.", "doc_access": "permission-blocked on experiments MCP"}
1778005052	2026-05-05	Zelun Wang	HF CLR ARF	https://docs.google.com/document/d/1tmnuZ1RjicqRKx8yBzOxK2wCvRHhDUA0aguizSAvWDI	https://pinterest.slack.com/archives/C05UMECTDDJ/p1778005051946709	Stabilizes CLR ARF after repeated reversions and bumps training dates one month forward. Strong sitewide repin and fresh-content gains, with a shopping clickout regression and a small Homefeed impression dip.	All Users					+0.81%	-0.41%	+0.45%												SSv2 Proxy +0.15% (non-SS, excluded from column); 5s+ Outbound Click +0.45%; Low Quality Sessions -0.34%; Fresh Neardup Impressions +2.49%; Fresh Neardup Repins +5.40%; [Homefeed] Impressions -0.26%; Closeups +0.38%; Repins +0.90%; Downloads +0.67%; Share +1.58%; Clickthrough +1.43%; Hides -4.71%	slack (doc blocked)	{"dropped_non_sig": {"SSv2 Proxy": "+0.15% explicitly marked non-SS by the poster"}, "notes": "SSv2 shopping clickout -0.41% is a genuine stat-sig regression on a canonical goal metric, recorded as such. Post is candid that ARF \"has gone through numerous reversions due to different failures\" and that the goal is a higher workflow success rate before handing to oncall - reliability is the real subject. Fresh Neardup Repins +5.40% is the largest fresh gain in the tracker.", "doc_access": "permission-blocked on experiments MCP"}
1777919860	2026-05-04	Zili Li	Improving LWS with Rank Distillation Loss	https://docs.google.com/document/d/1wgTVSE_Q0k_y_K_GavgyR92mvglsEuWtOzi7L9z8uRM/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1777919860405979	Introduces a rank distillation loss so the LWS student matches the teacher ranking order over items rather than just pointwise predictions, improving top-of-list fidelity. Modest SSv2 gains with solid Homefeed click lifts.	All Users	+0.06%	+0.37%																	SSv2 intent expressions only +0.20% (pre-AA, excluded from column); [Homefeed] Num Pin Click +1.15%; Clickthroughs +0.96%; Long Clickthroughs +1.06%	slack (doc blocked)	{"dropped_non_sig": {"SSv2 (intent expressions only)": "+0.20% flagged \"(pre-AA)\" by the poster, meaning the effect appears in the pre-experiment period and is likely not causal - kept out of the column"}, "notes": "The pre-AA annotation on intent expressions is the poster being appropriately careful; a reader skimming would likely have banked it as a win.", "doc_access": "permission-blocked on experiments MCP"}
1777065084	2026-04-24	Andreanne Lemay (She Her)	Retentive Recs: UIC diversification in SSD (V2)	https://docs.google.com/document/d/11nB5m6McocBa1ZntLTdtv_9YhDrL533xsV4Dv7goZgs/edit?usp=sharing	https://pinterest.slack.com/archives/C05UMECTDDJ/p1777065083952529	Integrates the user interest cluster V2 signal into SSD to diversify interest representation on Homefeed. V2 re-run with different cosine similarity thresholds and weight after V1 results proved irreproducible - V1 had run during the January fresh content spike, which inflated the treatment group.	All Users		+0.27%																	Closeup timespent (s) +0.22%; HF Repins +0.73%; Long clickthroughs +0.89%; Time spent (s) +0.21%; HF Pin Reports -14.82%; HF Pin Hides -2.61% prop; Num impressed Pincepts +0.17%; Num explored Pincepts +0.19%; TRADE-OFFS Pins 28d Fresh Neardup Impressions without Ads -0.68%; Pins 28d Fresh Neardup Repins without Ads -2.09%; Torso TP Non-Explored Impressions (All, Organic) -0.42% prop; Short clicks (non-product, non-ads) +0.96% (doc only); Home PWT cold start ~5ms and warm start ~5ms (doc only); model p90 latency 55ms to 60ms (doc only)	slack+doc	{"ship_group": "enabled_climit_20_ssd_uic_0x3_cs_0x8 (20 cluster limit, 0.8 cosine similarity threshold)", "cost": "Infra cost report states savings, but the doc says no significant infra change is expected", "notes": "METHODOLOGICALLY THE MOST HONEST LR IN THE TRACKER: the team withdrew V1 because the January fresh-content spike inflated its treatment group and results did not reproduce once the incident was fixed. Doc adds three things the post omits - a +0.96% non-product short-click regression, ~5ms Home PWT increase, and the model p90 going 55ms to 60ms. Doc also explains the fresh regression mechanistically (higher empty embeddings for fresh pins, lower cosine similarity between fresh pins and UICs) and notes Omnisage V2 should reduce it. Doc Learnings section records the tuning sweep: UIC term weight 1.0 caused stat-sig total SSv2 regression, 0.5 chosen; 20 clusters beat the 10 other workstreams use. This is a threaded reply to an older post (thread_ts 1770669768.287519), so the API returned it despite the top-level-only assumption.", "related_docs": ["https://docs.google.com/document/d/1JZmtHkk6M3VYEo9UMUUH9XEB-NOwABZIFS1njx7VDI ([Retentive SSD] Personalized Use Case Diversification)"]}
1776701720	2026-04-20	Hanlin Lu	Manas migration for Homefeed RecGPT	https://docs.google.com/document/d/1br2kXavSg9vJJBTY7UNnLelcHKlSuN4BWZic_zd52g4/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1776701719670599	Migrates HF-RecGPT ANN serving from Faiss to Manas. Neutral Helium metrics with $59,291/year savings and dev-velocity gains, unblocking the future SMS migration for the RecGPT pipeline.	All Users																			successful sessions v2 (total) +0.05% (non-stat-sig); Homefeed Repins -0.04% (non-stat-sig); Homefeed Pin Impressions +0.18% (non-stat-sig)	slack (doc blocked)	{"cost": "$59,291/year saving", "dropped_non_sig": {"SSv2 (total)": "+0.05% explicitly marked non-stat-sig by the poster"}, "notes": "Parity migration - all 18 fixed columns blank since every value is explicitly non-stat-sig. Earliest of this author's three HF-RecGPT LRs in the tracker (this, diversity improvements 2026-05-06, sizer reallocation 2026-06-30). The \"neutral cost on the newly approved Manas cluster\" claim in the May LR depends on this migration.", "doc_access": "permission-blocked on experiments MCP"}
1776363550	2026-04-16	Sophia Zhu	Feature Deprecation of indyRank and pageRank for HF CLR	https://docs.google.com/document/d/1Jfz5STnF6gKE9bvC2i7mtYE0n667ZFa7FZAdDyP0mg/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1776363549969239	Deprecates two stale viewer-side features (pinnability_viewer_pageRank and pinnability_viewer_indyRank) from the Homefeed CLR training dataset, retiring the workflows that generate them and reducing oncall load. Removing them also produced engagement gains.	All Users					+0.43%														5s+ Outbound Clicks +0.43%; [Homefeed] Clickthroughs +0.41%; Long Clickthroughs +1.22%; Picked For You Clickthroughs +0.89%; Interest Grid Clicks (Closeups) +1.15%; RED FLAG NUM_IMPRESSIONS with Low Quality Content Type Dimension -0.29%	slack (doc blocked)	{"notes": "Deleting two poorly-maintained features IMPROVED metrics - a clean argument for feature hygiene as engagement work, not just cost work. The flagged red flag (-0.29% low-quality-content impressions) is actually a desirable direction, which the post labels conservatively. Post links the exact feature-map source lines.", "doc_access": "permission-blocked on experiments MCP"}
1775853312	2026-04-10	Yali Bian	hf_shopping_cgs_to_lws_l1_utility	https://docs.google.com/document/d/1k1L2aQy-VHLMYTbZ3XHHa-gf9_JKrFgl9ADtYimyes/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1775853312102959	Onboards the remaining three Shopping L2R CG's to the unified LWS + L1 Utility pipeline, raising LWS/L1 coverage of CG budgets from ~75% to ~79%. Shopping CGs lose impression share but product-pin engagement rates improve sharply, evidence that LWS selects genuinely better product pins.	All Users					+0.32%														[Homefeed] Num Pin Click +0.22%; Repins +0.49%; Long Clickthroughs +0.84%; GPS neutral; Shop SRSD Onlyer +0.25% vol / +0.09% prop; Trustworthy Product 5s+ outbound click rate (homefeed, non-ad) +10.39%; TP Repin Rate (home feed) +3.23%; TP Closeup Rate (home feed) +1.13%; TP Longclick Rate (homefeed) +6.32%; Homefeed Closeup Rate +0.29%; Homefeed Long Clickthrough Rate +1.70%; LWS/L1 coverage of CG budgets ~75% to ~79%	slack (doc blocked)	{"notes": "MOST DISCUSSED LR IN THE TRACKER - 55 thread replies. This is the LWS shopping migration that the 2026-07-21 hf_shopping_control_l1_utility LR was later built to mitigate regressions from, so the two are one arc. All Trustworthy Product figures are RATES on the Homefeed surface, not sitewide counts - the +10.39% outbound click rate is the headline but it is a rate on a shrinking impression base, which is exactly the coverage-vs-quality distinction that matters here. Post is explicit that Shopping CGs LOSE impression share while engagement rates rise.", "doc_access": "permission-blocked on experiments MCP"}
1775758164	2026-04-09	Alok Malik	homefeed_dynamic_sizer_migration	https://docs.google.com/document/d/1xCez_Hy5aB22yk295yFatfedDPwXmdJ3UJRtmlTuD3w/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1775758163913739	Migrates Homefeed sizer computation from synchronous plan-time to asynchronous node-based resolve-time so model-driven sizing can run in parallel with resource fetching. Infra-only change that unblocks all future PHP models, saving ~40ms versus the static sizer approach.	All Users																			No major metric movement (infra-only change); ~20ms average latency added to homefeed, ~40ms saved vs the static sizer implementation; p90 PWT +25ms (doc Trade-offs)	slack+doc	{"ship_group": "enabled_node_model", "cost": "Neutral - doc notes the small increase in the cost report is probably noise, as the Manas clusters involved belong to other surfaces", "conflicts": {"PWT": {"slack": "\"PWT on homefeed increase by ~20ms, instead of 60ms\"", "doc": "Trade-offs section states p90 PWT +25ms"}}, "notes": "67 THREAD REPLIES - second most discussed in the tracker. TWO THINGS THE POST UNDERPLAYS: (1) the doc states this experiment CANNOT SHIP INDEPENDENTLY - the dynamic sizer and the PHP shopping load tuning model are co-dependent and both LRs must be approved before either ships. The Slack post does not mention that gating at all. (2) The post's \"~20ms instead of 60ms\" reads as a saving; the doc Trade-offs section separately records p90 PWT +25ms. The framing is net-positive-versus-alternative, not absolute-neutral. Doc also asks reviewers to compare control vs enabled_node_prod to confirm the migration alone changes no logic. James is a listed experiment design reviewer; Dylan is launch reviewer.", "related_docs": ["https://docs.google.com/presentation/d/1GGMnER8Gg5f1OWzVeVLZ-AyZPJg9XTaYvY57K9Tivzl (Dynamic Triggering: Personalizing CG Hyperparameters)", "https://docs.google.com/presentation/d/1UvBEKb5C1N6X2V9mp9qln8jeTuspfp_NhVJFAIGau0g ([ELT Presentation 3/30/26]: Learned Dynamic Triggering)"]}
1775667520	2026-04-08	Devin Kreuzer	BMI Relevance Filtering: Max Rate Tuning	https://docs.google.com/document/d/1Rep97Z8Z9iuoLoE3q5DTFX9dbUROG4qA7mAtJ5DoPQ/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1775667519866889	Raises the BMI relevance max filter rate from 50% to 100% and filters out pins with missing relevance scores instead of letting them bypass, to rule out causes of highly irrelevant feeds. Engagement largely neutral.	All Users																			Engagement metrics largely neutral - no values given	slack (doc blocked)	{"dropped_non_sig": {"all metrics": "post states engagement is largely neutral, no values given"}, "notes": "Shortest LR in April. Diagnostic in intent - the goal is ruling out causes of irrelevant BMI feeds, not moving metrics, so neutrality is the acceptable result. All 18 fixed columns blank. Note the mechanism: pins with MISSING relevance scores were previously bypassing the filter entirely.", "related_docs": ["https://docs.google.com/document/d/1EUYGKXr-eeveVDAEhLyzKylTivrrlUUTn_3UN5gHGuE (BMI relevance challenges doc)"], "doc_access": "permission-blocked on experiments MCP"}
1775519133	2026-04-06	Sujan Reddy Ale	HF CLR Router Simplification	https://docs.google.com/document/d/13HFtQOkAqiyd83FHpnf3_wA5yJpMwtGQQd_bqerQm8/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1775519133081589	Replaces per-condition bespoke CLR router features with a slot-based, condition-agnostic representation of two slot features per condition, plus a search sage-25 interest text embedding upgrade and DHEN feature cross weight tuning. Sitewide repin and Homefeed gains with two flagged red flags.	All Users							+0.39%												[Homefeed] Impressions +0.39%; Repins +0.83%; Clickthroughs +0.69%; Downloads on HF +0.91%; Pin screenshots +0.92%; Pin see more +2.97%; Closeups -0.55%; RED FLAGS Unsuccessful session +0.42% (inverted metric - positive is bad); Num impressions with low quality content type dimension +0.35%	slack (doc blocked)	{"notes": "Unsuccessful session +0.42% is an INVERTED metric where positive is bad, kept out of all SSv2 columns per the traps list, and paired with +0.35% low-quality-content impressions these are genuine quality red flags the poster labels as such. HF Closeups -0.55% while sitewide Repins +0.39% - mixed within-surface picture. Foundation win is real: adding a condition feature becomes a config mapping rather than a new feature set, unblocking Search CLR on HF and Semantic ID as a Condition. Same author as the 2026-06-02 UPP FM in CLR LR.", "related_docs": ["https://docs.google.com/document/d/1Gp5fhUKJMPuw0hRPtHVMIZuuLFZo7FqYuc_EXvsVWGo (Search CLR on HF)", "https://docs.google.com/document/d/1LGsABBa_50BoApIF-QC7Wbc1ooCGnlwzVcMuw6p_o5c (Semantic ID as a Condition)"], "doc_access": "permission-blocked on experiments MCP"}
1774905508	2026-03-30	Leo Zhang (He Him)	hf_multi_emb_funnel_efficiency	https://docs.google.com/document/d/146fJOaJl0NzHwm1_OMVERqBwFI72JoFhgd1KlRSSuno/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1774905507502699	Aligns the HF Multi Embedding CG with downstream ranking and blending using newly logged unimpressed pins data, adding position-slot information and a pairwise loss so the retrieval model learns full-funnel rather than recall-only signals. Repin and closeup gains against sizeable clickthrough regressions.	All Users		+0.185% vol / +0.05% prop					+0.49%	+0.35% vol / +0.05% prop		+0.17%									SSv2 (total) UCAN +0.11% prop; Unsuccessful Session v2 (total) UCAN -0.62% vol (inverted metric - negative is good); SSv2 Proxy UCAN +0.43% vol / +0.12% prop; Engaged Session +0.06% prop; UCAN Repin +1.08% vol; Homefeed Timespent +0.48% vol; [Homefeed] HF Repins +1.01% vol / +0.33% prop; HF Repins UCAN +1.43% vol / +0.38% prop; HF Closeup +0.69% vol / +0.09% prop; HF Closeup UCAN +1.13% vol; HF Impressions +0.17% vol; TRADE-OFFS Long Outbound Click -0.83% vol / -0.36% prop; Outbound Click -0.58% vol; HF Clickthrough -1.47% vol / -0.58% prop; HF Long Clickthrough -2.77% vol / -1.55% prop	slack (doc blocked)	{"notes": "HF Long Clickthrough -2.77% vol / -1.55% prop is one of the larger engagement regressions in the tracker and is disclosed plainly alongside the repin gains - a genuine repin-vs-click tradeoff, not a rounding artifact. \"SSv2 Proxy UCAN Vol +0.43\" is missing its % in the post; recorded verbatim in other_metrics rather than assuming. Full-stack change across data, model and unity-hf serving. Cross-team: UU, Homefeed and Infra.", "doc_access": "permission-blocked on experiments MCP"}
1774896001	2026-03-30	Devin Kreuzer	CLR GPU Serving via Attached Data	https://docs.google.com/document/d/1DlaW17dNx4KfKcBLamfQ1y-6JoKZ-Qz4WLLj43JnCxI/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1774896000652059	Migrates CLR to GPU Serving and restructures the CLR serving path, delivering $650k savings and an 85% model latency reduction. Metric picture is mixed - positive SRSD but neutral total SSv2, with a hypothesised CG distribution shift the team had not yet confirmed when seeking approval.	All Users			+0.13%	-0.14%	-0.40%		+0.40%												Low Quality Sessions -0.34%; [Homefeed] Impressions +0.26%; Repins +0.62%; Downloads +1.26%; Closeups -0.46%; Clickthroughs -0.57%; Long clickthroughs -0.85%; 85% reduction in model latency; 0.4% drop in HF e2e latency	slack (doc blocked)	{"cost": "$650k in cost savings from migrating to the new GPU cluster", "notes": "TWO GENUINE STAT-SIG REGRESSIONS ON CANONICAL GOAL METRICS: intent expressions -0.14% and non-shopping clickout -0.40%. The post is explicit that it is \"seeking LR approval with the current metric reading\" while a parallel experiment to confirm the CG-distribution-shift hypothesis and rule out training instability had only started that day - i.e. approval requested ahead of the diagnostic. Team hypothesis is that CLR CGs win distribution because of faster latency and fewer timeouts, which would make the tradeoff an artifact of the serving change rather than model quality. This is the GPU Serving launch that the 2026-05-29 CLR Condition Budget Tuning LR later exploits for spare capacity.", "doc_access": "permission-blocked on experiments MCP"}
1774030093	2026-03-20	David Wang	L1 Utility Presort Cutoff Multiplier 1.3x to 1.7x	https://docs.google.com/document/d/1ZPjiAbxJWG7iA0MUJZN_zukQTaprO2mAuugVhyixNA/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1774030092822119	Raises the L1 Presort cutoff multiplier from 1.3x to 1.7x per feed source, letting each source contribute more candidates into the cross-feed-source L1 Utility merge and rerank. Repin, freshness and diversity gains with a small PWT improvement and a modest cost saving.	All Users			+0.07% users				+0.35% vol / +0.16% users												Fresh Impressions +0.46% actions; Fresh Repins +0.28% users; Fresh Long Clickthroughs -0.59% users; Trustworthy Impressions +0.52% actions; Trustworthy Repins +1.56% users; Neardup Impressions Without Ads +0.87% actions / +0.1% users; Neardup Repins Without Ads +0.31% users; HF Impressions +0.03% users; Num Adopted Pincepts +0.25% actions; Num Adopted New Pincepts +0.27% actions; cold start PWT -0.2%	slack (doc blocked)	{"cost": "-$6.7k including only stat-sig services (saving)", "notes": "A single-parameter tuning change (1.3x to 1.7x) delivering broad small gains - the cheapest win-per-effort LR in the tracker. Third in the L1 Utility arc: hf_l1_utility, then UIC & SID diversity control, then this. \"Num Adopted Pincepts +0.25% actions\" is an HF-scoped use-case metric, not the canonical Adopted Pincept Domains (Interests), so it stays in other_metrics. A later reply in this thread (2026-04-01, David Wang) confirms full ramp-up.", "related_docs": ["https://docs.google.com/document/d/1rWeWZeTYOF0r0Hp_RDpyMQZNnIYECROnhHOdNRO6yIM ([LR] hf_l1_utility)", "https://docs.google.com/document/d/1gE9I03Wj6SpPlm-ARBGTZcXkt0jvBxp3B0nWzev_LW4 ([LR] Retentive Recs: UIC & SID for Diversity Control on L1 Utility)", "https://docs.google.com/document/d/1n38dNcDtblglQ7RC9JPNoeFzB2jnqq3QaKA-SNqjStE ([2026] L1 Utility on Homefeed roadmap)"], "doc_access": "permission-blocked on experiments MCP"}
1773963092	2026-03-19	Hedi Xia	Aligning Training with Serving in Homefeed Light Weight Scoring	https://docs.google.com/document/d/1CstdMibmwK0ci-JjyR0XqUSPm4_H71kPEh70DO3VbR0/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1773963091524149	Redesigns the LWS training objective from first principles with a new unimpressed dataset, trainer-side union of impressed and unimpressed data, and a pairwise distillation loss matching the LWS serving objective. Fixes a years-old training-serving target discrepancy and corrects a badly distorted closeup serving weight.	All Users	+0.10%			+0.15%				+0.19%											[Homefeed] Closeup +1.12%; Clickthroughs +0.77%; Hides -4.40%; Downloads +1.08%; [Core Users] Total SSv2 +0.25%; Intent expression +0.73%; Non shopping clickout +0.64%; Closeup +0.38%; 5s+ Outbound Clicks +0.61%; Board Creates -1.44%; HF Closeup +2.13%; HF Share -2.08%	slack+doc	{"ship_group": "lws_unimp_v4_20_6", "cost": "+$3,600/year net. Doc breaks this down: logging $114,646 + workflow $14,469 + 50% subsampled unimpressed S3 $286,390 + LWS prod iceberg $664,394 offset against LWS prod $1,076,300. Future deprecation of LWS prod data would unlock a further $664,394/year.", "notes": "THE MOST CONSEQUENTIAL FINDING IN THIS LR IS A RATIO, NOT A METRIC. The doc shows the repin:closeup weight ratio was 20:1 in training but 400:1 at serving - a 20x distortion that had been compensated for by artificially suppressing closeup weight for years. After this change serving sits at 16.67:1, close to training. The post alludes to \"a healthier and better-justified closeup serving weight\"; the doc quantifies it. Doc-only additions: Closeups +0.19% sitewide, Core Users non-shopping clickout +0.64% and 5s+ outbound +0.61%, plus two Core-User REGRESSIONS the post omits entirely (Board Creates -1.44%, HF Share -2.08%). Parent of the 2026-05-22 LWS Funnel Log Removal LR, which the doc pre-announces as a Q2 plan.", "related_docs": ["The Fundamental Theorem of Preranking (background theorem, linked in doc)"]}
1773417053	2026-03-13	Devin Kreuzer	Adding Pin CLR to BMI on GULP	https://docs.google.com/document/d/1ZGdfHZTXKlahltpDJwegHmWlW6kINbyqLETP03r6zzl/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1773417052639199	Adds Pin CLR to BMI on GULP, with large Board More Ideas surface gains and small sitewide propensity lifts. The GULP migration also unlocks Pin CLR for Notif Landing Page, Exploration Modules and AMB.	All Users																			Total SSv2 +0.03% prop; SSv2 Proxy +0.02% prop; Repins +0.18% (pval 0.102 - NOT stat-sig); US Market Repins +0.37%; [Board More Ideas] Impressions +0.81%; Repins +4.24%; Closeups +1.36%; Long Clickthroughs +2.66%; P2P Impressions +1.50%; Hides -7.50%; Trusted Product Repin/Clicks +15.90%; Timespent +0.45%	slack (doc blocked)	{"dropped_non_sig": {"Repins": "+0.18% at pval 0.102 - above the 0.05 threshold, so deliberately NOT placed in the Repins column despite the post listing it under Engagement Wins", "SSv2 (total)": "+0.03% propensity-only, reported without significance marking"}, "notes": "DATA-QUALITY CAVEAT DISCLOSED BY THE POSTER: an RTC logging issue on GULP prevented subtype and visualization reports from running at all. The team committed to fixing it before rollout and rerunning the analyses in the CLR holdout - so the subtype evidence behind this LR did not exist at review time. All 18 fixed columns blank: the sitewide numbers are propensity-only or non-significant, and everything else is BMI-surface scoped.", "related_docs": ["https://docs.google.com/document/d/1g_QeAdatGYwR1fbgF26E59gfuo1_-65j80xnfjwCNJU ([Notif LR] Context-Aware Recommendations for Notif Landing Page)"], "doc_access": "permission-blocked on experiments MCP"}
1773177169	2026-03-10	Hedi Xia	Homefeed Light Weight Scoring Unified Tower and Transact	https://docs.google.com/document/d/18rWiT5dbiv1GFDSv7BheXYoLT3VFG5xtwmdAWawoBGE/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1773177168614229	Replaces the split CPU-viewer-tower plus GPU-pin-tower LWS architecture with a Unified Tower on GPU and onboards the Transact module, which was not viable under two towers. Large hide and report reductions at no added serving cost, unblocking foundation models and 16k sequences.	All Users	+0.10%		+0.12%	+0.10%															SSv2 non-shopping clickout MECE +0.03%; SSv2 shopping clickout -0.04%; [Homefeed] Hides -4.22%; Reports -12.72%; Screenshot +0.85%; Closeup (s) +0.16% unique user; Clickthroughs -0.92%; Long Clickthroughs -0.99%	slack+doc	{"ship_group": "v8p7p1_nogate_unityeff_3_gpu (Unified Tower GPU, Transact module, 32 most-recent items)", "cost": "No added serving cost per the post; doc has a dedicated Cost section", "notes": "Doc adds two canonical values absent from the post: SSv2 intent expressions +0.10% MECE and shopping clickout -0.04%. Doc explains the architectural payoff clearly: the two-tower design mixed and compressed user features early, obscuring individual signal contributions; the unified tower preserves demographic features (which drive impressions) separately from tokenized sequence features (which predict repins). Doc also documents the rejected alternative (16 tokens instead of 32) as having more balanced HF metrics and stronger HF repin - i.e. the shipped group was NOT the best on HF repin, chosen for intent-expression sessions instead. Reports -12.72% is the largest report reduction in the tracker. Same author as the 2026-03-19 LWS alignment LR nine days later.", "related_docs": ["https://docs.google.com/document/d/1WHknCtOCQ-T2JQvsXyq77hgQSTg4OlfhoVLW2nL11_0 (HF LWS Unified Tower)", "https://docs.google.com/document/d/171lWLwPPQ54bHrxCKvrPuf8mBirq28kQOLqgCMzpa-o (LWS: GPU Serving)"]}
1772557075	2026-03-03	Chuxi Wang	Retentive Recommendation: Frontier Sampling in UIC x CLR	https://docs.google.com/document/d/1Dn3dNzZcG-Uuc441l-DbyEoWc9H34k7HmqWRmaBOuu8/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1772557075488349	Introduces frontier sampling to select UIC landmarks at the boundary of a user known interest clusters, shifting UIC CLR content toward a more exploratory distribution. Broad topline and Homefeed gains with confirmed diversity increases, plus early evidence for predicted UIC.	All Users		+0.22%	+0.19%				+0.58%			+0.24%	+0.35%								Board Creates +0.67%; [Homefeed] HF Pin Impressions +0.48%; Num Pin Click +0.53%; Repins +0.80%; Downloads on HF +1.88% vol / +0.83% unique users; Pin Screenshots from HF +2.06% vol / +1.06% unique users; Long Clickthroughs -0.55% unique users; Num Impressed Pincepts +0.22%; Num Explored Pincepts +0.28%; Num Impressed New Pincepts +0.19%; Num Explored New Pincepts +0.32%; Home Feed timespent (s) +0.41%; Homefeed Use Case Adoption +0.38%; Impression Diversity +0.25%	slack (doc blocked)	{"ship_group": "score_sort_sel_5_ft_th_90", "notes": "Earliest UIC x CLR LR in the tracker and the origin of the predicted-UIC (pUIC) line of work that later drives the Pred UIC CLR budget architecture. Post states this experiment provides \"early evidence for the value of predicted UIC\". \"SSv2 SRSD +0.19%\" mapped to the MECE column on the cross-LR basis established by the 2026-06-16 Downstream Rewards doc. All Pincept metrics here are Num-prefixed HF-scoped counts, not the canonical Domains (Interests) metrics.", "doc_access": "permission-blocked on experiments MCP"}
1772135159	2026-02-26	Yidi Wang	Homefeed Multi-embedding Retrieval: Deprecation of Annotation V5	https://docs.google.com/document/d/16LfZ_oJmRUYY2sya0W8DCU6zKmUGtRl41ZteHfyFQNg/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1772135158583849	Removes the annotation v5 feature from Multi-Embedding retrieval. Reported metrics are slightly positive across the board, but the poster attributes this to pre-AA bias and characterises the true net effect as neutral.	All Users																			Total SSv2 +0.11%; SSv2 Proxy +0.16%; Intent Expression MECE +0.20%; SRSD +0.06%; Non-Shopping Clickout +0.13%; Shopping Clickout +0.03%; Pin Impressions +0.14%; Closeup(s) +0.22%; Repins +0.22%	slack (doc blocked)	{"dropped_non_sig": {"all reported values": "poster states the slightly positive SSv2 is likely a non-stat-sig positive pre-AA bias and that the net effect of removing the feature is NEUTRAL, with no HF degradation once AA bias is excluded"}, "notes": "ALL 18 COLUMNS DELIBERATELY BLANK despite the post listing nine positive values. The poster explicitly disowns them as pre-AA artifact - recording them as wins would invert the author's own conclusion. Companion to the 2026-01-12 LWS annotation v5 deprecation by the same author (Yidi Wang): same feature retired from two different models seven weeks apart.", "doc_access": "permission-blocked on experiments MCP"}
1771016200	2026-02-13	Raymond Hsu	See More on Pins CLR	https://docs.google.com/document/d/1dwlpbMWIACUMAJDqNpJbLF-EywONo4D5LxFLefXGpE/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1771016199858609	Adds See More actions to the Pins CLR query composer to maintain See More coverage after Pins CLR replaced CGs that used it (Coengagement, part of IPFY), and to keep action types consistent across pin-to-pin CGs. Mostly neutral with stat-sig closeup lifts on Search and Related Pins.	All Users																			[stat-sig] RP Closeups +0.22%; Search Closeups +0.21%; [non-stat-sig] Intent Expression -0.02%; SSv2 +0.05%; SRSD +0.11%; Impressions +0.14%; Closeups +0.14%; repins +0.07%; 5s+ clickthroughs +0.13%; Hides +0.68%; Reports +6.74%; HF Impressions +0.10%; HF Closeups +0.09%; HF Repins -0.05%; HF Clickthroughs +0.05%; HF Hides +2.25%; HF Reports +9.55%	slack (doc blocked)	{"dropped_non_sig": {"all SSv2 and sitewide values": "poster marks stat-sig with bold+underline; only RP Closeups and Search Closeups qualified. HF Reports +9.55% and Hides +2.25% are non-stat-sig and attributed to a pre-AA effect"}, "notes": "Exemplary significance labelling - the poster states the convention explicitly (\"bold and underline = stat-sig, otherwise non-stat-sig\"), which is why all 18 columns are blank: the only two stat-sig values are Search/RP-surface closeups. The HF Reports +9.55% would look alarming stripped of that caveat.", "doc_access": "permission-blocked on experiments MCP"}
1770669768	2026-02-09	Andreanne Lemay (She Her)	Retentive Recs: UIC diversification in SSD (V1)	https://docs.google.com/document/d/11Kzu_QnQLRvdk6htg5OxjLxOgPKNMdq7LTyorGGukQ/edit?usp=sharing	https://pinterest.slack.com/archives/C05UMECTDDJ/p1770669768287519	Integrates the UIC V2 signal into SSD to diversify interest representation on Homefeed. Strong repin and retention gains at the time of posting - but this version was later WITHDRAWN as irreproducible.	All Users	+0.13%		+0.20%		+0.71%		+0.79%												HF Repins +1.74%; Sessions >5min +0.29%; Time spent (s) +0.24%; Num impressed Pincepts +0.17%; Num adopted Pincepts +0.23%; TRADEOFFS Pins 28d Fresh Impressions -1.67%; Pins 28d Fresh Repins -1.67%; Short clicks (non-product, non-ads) +0.85%; +5ms SSD p90 latency (from ~60ms)	slack (doc blocked)	{"notes": "*** THIS LR WAS WITHDRAWN. *** Per the V2 post on 2026-04-24 (same author, threaded under THIS message), the team did not launch this version: the experiment ran during the January fresh-content spike, which inflated the treatment group, and the results were NOT REPRODUCIBLE once that incident was resolved. The values in the fixed columns are therefore what was ASSERTED on 2026-02-09, not what held up - consistent with the tracker being an account of claims at review time, but this row should never be cited as a result. V2 (2026-04-24) with different cosine thresholds and weight is the version that stands. Post typo \"Non shipping click out\" read as non-shopping clickout.", "related_docs": ["https://docs.google.com/document/d/11nB5m6McocBa1ZntLTdtv_9YhDrL533xsV4Dv7goZgs ([LR] Retentive Recs: UIC diversification in SSD V2 - the replacement)"], "doc_access": "permission-blocked on experiments MCP"}
1769641977	2026-01-28	Hanyu Li	BMI RecGPT Unified Retrieval/Ranking CG	https://docs.google.com/document/d/19ht7cFrwi4FA9ZkpwBoXWtZfEJcwgevLAo9ZUj3xug/edit?usp=sharing	https://pinterest.slack.com/archives/C05UMECTDDJ/p1769641978895319	Brings RecGPT to the BMI surface for the first time with a board-prompted variant balancing relevance and engagement, and a unified L0 retrieval plus L1 lightweight scoring service in one ensemble. Large BMI repin gains at ~$295k/yr cost.	All Users							+0.29% vol / +0.11% users	+0.01% vol / +0.03% users											SSv2 Proxy +0.01% vol / +0.03% users; SSv2 Save Sessions +0.11% vol / +0.10% users; [BMI] More Ideas Num Repin +3.63% vol / +2.82% users; More Ideas Board Session +0.33% vol / +0.32% users; More Ideas Num Impression -0.25% vol / +0.32% users; More Ideas Num Pin Clicks +0.42% vol / +1.72% users; More Ideas Num Outbound Clicks +0.56% vol / +0.62% users; Num Impressed Pincepts +0.08% / +0.04%; Num Adopted Pincepts +0.18% vol / +0.10% users	slack (doc blocked)	{"cost": "+~$295k/yr ($177k serving RecGPT + $118k impact on L2 ranking), with further savings expected from the upcoming GULP migration and WIP model iterations", "notes": "First demonstration of unified L0 + L1 serving in one ensemble via reuse of the input user sequence and ANN fetch results - the architectural claim matters more than the metrics. Sitewide Repins +0.29% and Closeups +0.01% are technically sitewide so they fill those columns, but note the closeup value is at the noise floor. This is the LR whose follow-up (2026-03-27 thread reply) revalidated results after the unity-hf to unity-gulp migration and secured $123k + $116k budget approvals.", "doc_access": "permission-blocked on experiments MCP"}
1769480100	2026-01-27	Xinyi Li	HF RecGPT CG	https://docs.google.com/document/d/1LE_oe2s5qDXcX9siKlxabc01BVRgM4e3Xbkmm6PaCBU/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1769480099905659	Establishes RecGPT as a new Homefeed candidate generator after improvements across data, modeling and infra to the autoregressive model. Strong engagement with serving cost cut from $970k to ~$571k/yr versus the previous launch candidate.	All Users			+0.10% users				+0.37% users												Shares +0.84% vol; [Homefeed] Repin +0.71% vol / +0.86% users; Closeup +0.16% users; Impression -0.29% vol; Shares +2.89% vol / +2.17% users; Num Impressed Pincepts +0.16% vol; Num Adopted Pincepts +0.35% vol / +0.16% users	slack (doc blocked)	{"cost": "+~$571k/yr serving + ranking, down from $970k for the previous launch candidate; further savings expected from upcoming infra migration", "notes": "15 thread replies. Origin of RecGPT as an HF candidate generator - the 2026-04-20 Manas migration, 2026-05-06 diversity improvements and 2026-06-30 sizer reallocation all descend from this. Post pre-empts the impression drop: HF Impression -0.29% is attributed to an engagement shift toward higher positions and explicitly \"not considered a loss\". Cost is the story - $571k is still large, and a later thread reply (2026-03-12, Xinyi Li) confirms cost was the blocker and was brought down to ~$300k by replacing 100 CLR pins.", "doc_access": "permission-blocked on experiments MCP"}
1769104749	2026-01-22	David Wang	Retentive Recs: UIC & SID for Diversity Control on L1 Utility	https://docs.google.com/document/d/1gE9I03Wj6SpPlm-ARBGTZcXkt0jvBxp3B0nWzev_LW4/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1769104749291599	First application of the L1 Utility layer for global cross-CG control, using UIC and semantic ID assignments as diversifying signals. Impression diversity gains drive SSv2, retention and use-case wins with minimal PWT and latency cost.	All Users	+0.10% vol / +0.07% users		+0.08% users	+0.13% vol / +0.08% users								+0.06%							Sessions >5m +0.30% vol / +0.11% users; Fresh Impressions +0.19% users; Trustworthy Pins Fresh Impressions +0.34% users; Fresh Neardup Impressions Without Ads +0.14% users; [Homefeed] Impressions +0.28% actions; Repins +0.19% users; Long Clickthrough +0.44% users; Reports -4.05% users; Reactions +0.48%; Num Impressed Pincepts +0.16%; Num Explored Pincepts +0.20%; Num Adopted Pincepts +0.18%; Num Impressed New Pincepts +0.16%; Num Explored New Pincepts +0.24%; Num Adopted New Pincepts +0.23% actions / +0.13% users	slack (doc blocked)	{"cost": "+$37k, attributed to search wins - the increase comes from manas-searchsage_embeddings-baseprimeexp_v2, likely more users searching after improved diversity", "notes": "SECOND in the L1 Utility arc (after hf_l1_utility, before the 2026-03-20 presort multiplier and 2026-06-11 Iterative Diversity). Notable cost framing: the $37k increase is a CONSEQUENCE OF SUCCESS - better diversity drove more searching, which cost more in search embeddings. That is a rare and honest way to report a cost line.", "doc_access": "permission-blocked on experiments MCP"}
1769031767	2026-01-21	Sophia Zhu	Reweighting DHEN Feature Inputs for Homefeed CLR	https://docs.google.com/document/d/1xUf0e3joJnnEOQMIsYaKw321-ZYuf_90rSdoPrKe-yU/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1769031767040289	Retunes the number of tokens learned per feature in the CLR DHEN feature-crossing layer, upweighting newly launched transformer-encoded sequence and condition tokens while downweighting continuous features. Gains concentrated in CLR-powered CGs, especially Interest CLR.	All Users						+0.27%													[Homefeed] Repins +0.36%; Downloads on HF +0.53%; Use Case Retention +0.28%; Pin Share Completed +0.97%; Homefeed Interest Pin Impressions +0.76%; Interest Repins +1.32%; Interest Clickthroughs +1.09%; Interest RCs +1.31%; Promoted Pin Clickthroughs +0.55%; Promoted Pin Repins or Clickthroughs +0.57%; RED FLAG Pin Reports -7.88% (favourable direction); Related Reports -11.81%; [Growth overall] Repins UCAN +0.69%; [Low-Signal] SSv2 MECE +0.48% prop; [Higher-Engaged] Repins UCAN +0.66%	slack (doc blocked)	{"notes": "A pure hyperparameter retune (token counts per feature group) delivering broad gains - like the 2026-03-20 presort multiplier, evidence that tuning existing levers competes with new features. The two \"Red Flags\" the poster lists (Pin Reports -7.88%, Related Reports -11.81%) are both in the FAVOURABLE direction, so the label appears to be section convention rather than a warning. Same author as the 2026-04-16 indyRank/pageRank deprecation LR.", "doc_access": "permission-blocked on experiments MCP"}
1768924779	2026-01-20	Zelun Wang	Improve long-tail performance in CLR with Semantic IDs	https://docs.google.com/document/d/1QamRyk7X2mQbWTSO8uC_rg4aR922XQwy-I5Ez_a1H8c/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1768924779474889	Adopts Semantic IDs as a pin-side feature in CLR to improve long-tail performance, testing three featurization methods. Best group delivers session, retention and diversity gains, with effects strongest for new user cohorts.	All Users	+0.13%		+0.15%	+0.09% vol / +0.07% users				+0.16%		+0.14% vol / +0.05% users									Pins 28d fresh impressions +0.14% vol / +0.08% users; [Homefeed] HF impressions +0.16%; HF Clickthroughs +0.54% vol / +0.12% users; HF Screenshots +0.76% vol / +0.55% users; HF Interest Repins +1.76% vol / +1.00% users; DAU (Web Or API Requests unique users) +0.05%; UCAN MAU +97,426; Use Case Intention +0.29%; HF Use Case Adoption +0.47% vol / +0.12% users; Impression Diversity +0.17% vol / +0.06% users	slack (doc blocked)	{"notes": "DAU is reported as \"DAU (Web Or API Requests unique users) +0.05%\" - a proxy definition rather than the canonical DAU metric, so it stays in other_metrics rather than the DAU column. UCAN MAU +97,426 is an absolute count on a segment, so MAU stays blank. Companion to the 2026-02-10 PinCLIP Semantic ID in FM LR: same technique applied to retrieval (here) and ranking (that one), three weeks apart. Three featurization methods tested - the alternatives are the useful content in the doc.", "doc_access": "permission-blocked on experiments MCP"}
1768424328	2026-01-14	Hanlin Lu	Efficiency improvements in Multi Embedding Learned Retrieval	https://docs.google.com/document/d/1-WoYYkHdwOPs4aPYUQg2zT_NLuBySYZUTE8kARZe_DU/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1768424327996589	Combines training-infra improvements, hyperparameter tuning and system tuning to cut Multi Embedding model training time from 19.3 hours to 4.75 hours, a 4x speedup, while also delivering clickthrough gains.	All Users					+0.67% vol / +0.41% users														ME model training time 19.3h to 4.75h (4x speedup); [Homefeed] Pin Impressions +0.20% vol / +0.08% users; Num Pin Click +0.70% vol / +0.07% users; Clickthroughs +1.85% vol / +1.25% users; Long Clickthroughs +2.02% vol / +1.64% users; Pin Share Completed +1.47% vol / +0.76% users; TRADEOFFS Downloads on HF -0.48% vol / -0.19% users; Pin Screenshots from HF -0.73% vol / -0.43% users	slack (doc blocked)	{"notes": "14 thread replies. The 4x training speedup (19.3h to 4.75h) is the headline and a pure dev-velocity win; the engagement gains came along with it, which is unusual for an efficiency LR. Download and screenshot regressions disclosed. Same author as the RecGPT series (Manas migration, diversity improvements, sizer reallocation).", "doc_access": "permission-blocked on experiments MCP"}
1768258011	2026-01-12	Yidi Wang	Homefeed LWS Deprecate Annotation v5	https://docs.google.com/document/d/1Hiivr5z55ZvYMenaUVQ3WZ_dpik2q7WzEJoxdg8buQY/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1768258011288999	Removes the annotation v5 feature from LWS ahead of the feature being deprecated platform-wide. SSv2 neutral with small mixed movements; no incremental cost saving since it was already realized in GPU serving benchmarks.	All Users									+0.16%										SSv2 neutral; Closeup (s) +0.25%; [Homefeed] HF Closeup (s) +0.31%; HF Repin -0.49%; HF Clickthrough +0.52%; HF Long Click +0.89 (% sign missing in post)	slack (doc blocked)	{"dropped_non_sig": {"SSv2": "stated neutral, no values given"}, "notes": "\"Closeup (s)\" here is TIMESPENT in seconds, not the Closeups action count, so the Closeups column stays blank per the traps list - only Pin Searches +0.16% fills a canonical column. HF Repin -0.49% is a disclosed regression. \"HF Long Click: +0.89\" is missing its % sign in the post; recorded verbatim. Post is explicit that the cost saving was ALREADY realized in GPU serving benchmarks and this yields no further saving - a useful guard against double-counting. Companion to the same author's 2026-02-26 Multi-Embedding annotation v5 deprecation.", "doc_access": "permission-blocked on experiments MCP"}
1767991644	2026-01-09	Yali Bian	Onboard BoardCLR to LWS and L1 Utility	https://docs.google.com/document/d/1ce8PYvoUaQSbY1gChl015JkWJv16Ga41qCf_CiXxieY/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1767991644190129	Onboards BoardCLR to the unified LWS plus L1 Utility pipeline as a pilot for migrating all CGs, raising LWS/L1 coverage of CG budgets from 68.9% to 75.5%. Metrics neutral sitewide but BoardCLR own engagement rates improve sharply, with a small cost saving.	All Users																			SSv2 and Homefeed metrics neutral; BoardCLR repin rate +0.10% absolute (+16.3% relative); BoardCLR closeup rate +0.43% absolute (+11.3% relative); LWS/L1 coverage of CG budgets 68.9% to 75.5%; 27% duplicate rate found between BoardCLR and other CGs	slack (doc blocked)	{"cost": "~$21K/year savings", "dropped_non_sig": {"SSv2 and Homefeed": "stated neutral, no values given"}, "notes": "FIRST STEP of the CG-onboarding arc that the 2026-04-10 hf_shopping_cgs_to_lws LR continues (75% to 79% coverage) - this one moved 68.9% to 75.5%. All 18 columns blank: sitewide is neutral by design, and the wins are CG-internal rates. The +16.3% RELATIVE repin rate reads far larger than the +0.10% absolute - both are given, which is good practice. The 27% duplicate rate between BoardCLR and other CGs is the finding that justifies L1 Utility dedup.", "related_docs": ["https://docs.google.com/document/d/1NEqMnYaxbTDl_9BJ4A9TL0xo9kLhblMT3PShg_mHRu4 (Diversity Control on L1 Utility)"], "doc_access": "permission-blocked on experiments MCP"}
1767906151	2026-01-08	Hanlin Lu	Multi Embedding UOE LR follow-up (add UOE module)	https://docs.google.com/document/d/1KDr5wiOir8rJ7noJxuZXf567hds13uOO4IFr21jKLUA/edit?tab=t.0	https://pinterest.slack.com/archives/C05UMECTDDJ/p1767906151394509	Follow-up to an already-approved LR that added three UOE/DERM features to Multi Embedding: adding a new UOE module on top makes every metric more pronounced, so the team proposes shipping features plus module instead.	All Users	+0.12%			+0.15%															[3 features + UOE module, the proposed ship] Homefeed Num Pin Click +1.14% vol / +0.11% users; Related Impressions Diluted +0.23% vol / +0.03% users; Num Explored Pincepts +0.44% vol / +0.04% users; [3 features only, previously approved] Total SSv2 +0.08%; Intent Expressions MECE +0.10%; Homefeed Num Pin Click +0.91% vol / +0.14% users; Related Impressions Diluted +0.15%; Num Explored Pincepts +0.39% vol / +0.06% users	slack (doc blocked)	{"notes": "AMENDMENT TO A PRIOR APPROVAL, not a new LR - a threaded follow-up (thread_ts 1759273867.020149) asking to change the ship group after approval. Fixed columns hold the PROPOSED group (3 features + UOE module); the previously-approved 3-features-only figures are labelled in other_metrics so the two are not conflated. Team states they updated the original doc's Summary and Highlighted Metric Gains sections but deliberately left the rest unchanged to preserve the original rationale - good doc hygiene worth noting.", "doc_access": "permission-blocked on experiments MCP"}
1767729395	2026-01-06	Chuxi Wang	Retentive Recommendations: UIC Signals in CLR	https://docs.google.com/document/d/116-NuyhfFqxmYJgMkMOQui1g1yvBz-z-SYWv6I6RO8/edit?usp=sharing	https://pinterest.slack.com/archives/C05UMECTDDJ/p1767729395300429	Builds the User Interest Cluster V2 signal and integrates it fully into CLR, replacing legacy followed interest. Engagement, retention and diversity gains plus $322k savings, with UIC CG engagement rates on par with top ML candidate generators.	All Users	+0.11% users		+0.19% vol / +0.13% users		+0.41%								+0.17%		+0.14%				[Homefeed] HF impressions +0.09% users; HF Closeups -0.34% actions/user but +0.16% unique users; HF Repins +0.69%; HF Clickthroughs +0.79%; HF Long Clickthroughs +1.28%; HF Downloads +2.87%; HF Screenshots +2.02%; Pin Share Completed on HF +1.34%; Num Impressed Pincepts +0.39%; Num Adopted Pincepts +0.55%; Num Impressed New Pincepts +0.17%; Num Adopted New Pincepts +0.29%; Use Case Retention +0.42%; Pins 28d fresh impressions +0.10% users; Pins 28d fresh Neardup impressions without ads +0.49% vol / +0.15% users; Trustworthy Product Pin Impressions +0.69%	slack (doc blocked)	{"cost": "$322k saving - from deprecating followed interest and adding fewer candidates for the UIC x CLR CG", "notes": "FOUNDATION OF THE ENTIRE RETENTIVE RECS / UIC WORKSTREAM - first LR in the tracker chronologically for this line, and the post explicitly names the next phase: \"user interest exploration with predicted UICs\" (pUIC), which arrives as the 2026-03-03 Frontier Sampling LR. Deprecates legacy followed interest, which matches the known-deprecated status of Followed Interest RTC 5. Third row in the tracker to populate DAU and WAU as clean percentages. HF Closeups is negative in actions/user but positive in unique users - a reach-vs-intensity split worth noting.", "doc_access": "permission-blocked on experiments MCP"}